



King County Real Estate

1. Overview

This notebook examines King County, WA dataset of houses and reviews how and what renovations add value to a house's sale price.

2. Business Problem

The goal is to provide advice to homeowners about how home renovations might increase the estimated value of their homes, and by what amount.

3. Data Understanding

The data comes from the official King County WA website: kingcounty.gov and more specifically from the section King County Department of Assessments. More information can be found [here](https://info.kingcounty.gov/assessor/DataDownload/default.aspx%5D) (<https://info.kingcounty.gov/assessor/DataDownload/default.aspx%5D>).

This dataset contains information about house sales in King County WA. It includes the following columns:

- `id` - Unique identifier for a house
- `date` - Date house was sold
- `price` - Sale price. This criteria will be the prediction target.

- bedrooms - Number of bedrooms
- bathrooms - Number of bathrooms
- sqft_living - Square footage of living space in the house
- sqft_lot - Square footage of the lot
- floors - Number of floors in house
- waterfront - Indicates if the house is on a waterfront
- greenbelt - Indicates if the house is near to an area of open natural, undeveloped land
- nuisance - Indicates if the house has traffic noise or other nuisances
- view - Quality of view from house
- condition - Overall condition of the house
- grade - Overall grade of the house
- heat_source - Heat source for the house
- sewer_system - Sewer system for the house
- sqft_above - Square footage of house apart from basement
- sqft_basement - Square footage of the basement
- sqft_garage - Square footage of garage space
- sqft_patio - Square footage of outdoor porch or deck space
- yr_built - Year when house was built
- yr_renovated - Year when house was renovated
- address - Street address
- lat - Latitude coordinate
- long - Longitude coordinate

3. a- Loading the data using pandas

```
In [1]: ┏━━━ import pandas as pd
      ━━━ import numpy as np
      ━━━ import matplotlib.pyplot as plt
      ━━━ import matplotlib.ticker as ticker
      ━━ from matplotlib.patches import Rectangle
      ━━ import seaborn as sns
      ━━ import sqlite3
      ━━ import warnings
      ━━ import statsmodels.api as sm
      ━━ from sklearn.metrics import mean_absolute_error

      %matplotlib inline
```

```
In [2]: ┏━━━ kc_house_df = pd.read_csv("data/kc_house_data.csv")
```

4. Data Preparation

In [3]: ► kc_house_df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30155 entries, 0 to 30154
Data columns (total 25 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   id               30155 non-null   int64  
 1   date              30155 non-null   object  
 2   price              30155 non-null   float64 
 3   bedrooms            30155 non-null   int64  
 4   bathrooms            30155 non-null   float64 
 5   sqft_living          30155 non-null   int64  
 6   sqft_lot              30155 non-null   int64  
 7   floors              30155 non-null   float64 
 8   waterfront            30155 non-null   object  
 9   greenbelt             30155 non-null   object  
 10  nuisance              30155 non-null   object  
 11  view                 30155 non-null   object  
 12  condition             30155 non-null   object  
 13  grade                 30155 non-null   object  
 14  heat_source            30123 non-null   object  
 15  sewer_system            30141 non-null   object  
 16  sqft_above              30155 non-null   int64  
 17  sqft_basement            30155 non-null   int64  
 18  sqft_garage              30155 non-null   int64  
 19  sqft_patio              30155 non-null   int64  
 20  yr_built                30155 non-null   int64  
 21  yr_renovated             30155 non-null   int64  
 22  address                 30155 non-null   object  
 23  lat                     30155 non-null   float64 
 24  long                    30155 non-null   float64 
dtypes: float64(5), int64(10), object(10)
memory usage: 5.8+ MB
```

The dataset has:

- 25 columns: 10 integers, 5 float, 10 objects.
- 30155 rows. 'heat_source' is the field with the highest amount of null-values: 32. Dropping them allows to keep 99.9% of the data, they were dropped without deteriorating.

4. a- Dropping null values

In [4]: ► `# Dropping null values`
kc_house_df.dropna(inplace=True)

In [5]: # Verifying null values were dropped
kc_house_df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 30111 entries, 0 to 30154
Data columns (total 25 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   id                30111 non-null   int64  
 1   date              30111 non-null   object  
 2   price              30111 non-null   float64 
 3   bedrooms           30111 non-null   int64  
 4   bathrooms          30111 non-null   float64 
 5   sqft_living        30111 non-null   int64  
 6   sqft_lot            30111 non-null   int64  
 7   floors              30111 non-null   float64 
 8   waterfront          30111 non-null   object  
 9   greenbelt           30111 non-null   object  
 10  nuisance             30111 non-null   object  
 11  view               30111 non-null   object  
 12  condition           30111 non-null   object  
 13  grade               30111 non-null   object  
 14  heat_source         30111 non-null   object  
 15  sewer_system        30111 non-null   object  
 16  sqft_above           30111 non-null   int64  
 17  sqft_basement       30111 non-null   int64  
 18  sqft_garage          30111 non-null   int64  
 19  sqft_patio           30111 non-null   int64  
 20  yr_built             30111 non-null   int64  
 21  yr_renovated        30111 non-null   int64  
 22  address              30111 non-null   object  
 23  lat                  30111 non-null   float64 
 24  long                 30111 non-null   float64 
dtypes: float64(5), int64(10), object(10)
memory usage: 6.0+ MB
```

In [6]: # Inspecting statistics about the dataset
kc_house_df.describe()

Out[6]:

	id	price	bedrooms	bathrooms	sqft_living	sqft_lot
count	3.011100e+04	3.011100e+04	30111.000000	30111.000000	30111.000000	3.011100e+04
mean	4.539171e+09	1.108971e+06	3.415197	2.335708	2113.342798	1.664880e+04
std	2.882236e+09	8.965158e+05	0.979755	0.888293	973.453260	5.993303e+04
min	1.000055e+06	2.736000e+04	0.000000	0.000000	3.000000	4.020000e+02
25%	2.070850e+09	6.492360e+05	3.000000	2.000000	1420.000000	4.850000e+03
50%	3.876000e+09	8.600000e+05	3.000000	2.500000	1920.000000	7.477000e+03
75%	7.287100e+09	1.300000e+06	4.000000	3.000000	2620.000000	1.056800e+04
max	9.904000e+09	3.075000e+07	13.000000	10.500000	15360.000000	3.253932e+06

The mean price is \$1,108,971 for houses of over 3 bedrooms and 2 bathrooms, with a living space of 2,113 square feet on average and 1,810 square feet above the basement. The average year built is 1975. The dataset contains houses built from 1900 to 2022.

In [7]: # Inspecting the dataset
kc_house_df.head()

Out[7]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront
0	7399300360	5/24/2022	675000.0	4	1.0	1180	7140	1.0	
1	8910500230	12/13/2021	920000.0	5	2.5	2770	6703	1.0	
2	1180000275	9/29/2021	311000.0	6	2.0	2880	6156	1.0	
3	1604601802	12/14/2021	775000.0	3	3.0	2160	1400	2.0	
4	8562780790	8/24/2021	592500.0	2	2.0	1120	758	2.0	

5 rows × 25 columns

4. b- Removing unnecessary columns

The columns below are removed because they do not provide any value to the analysis to understand what renovations add value to a house.

Anything pertaining to the house location can be removed: will not have impact on renovations.

- lat
- long
- address
- greenbelt
- waterfront
- view

Anything pertaining to dataset admin can be removed

- id
- date (date house was sold)
- yr_renovated

```
In [8]: ┏ unnecessary_columns = ["lat", "long", "address", "greenbelt", "waterfront",  
In [9]: ┏ kc_house_df = kc_house_df.drop(columns=kc_house_df[unnecessary_columns])  
In [10]: ┏ # Verifying new columns  
         kc_house_df.columns  
  
Out[10]: Index(['price', 'bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors',  
                 'nuisance', 'condition', 'grade', 'heat_source', 'sewer_system',  
                 'sqft_above', 'sqft_basement', 'sqft_garage', 'sqft_patio', 'yr_built'],  
                 dtype='object')  
  
In [11]: ┏ len(kc_house_df.columns)  
  
Out[11]: 16
```

The dataset now has 16 columns

5. Simple Linear Regression Model: Initial Baseline

A simple linear regression model is a statistical method used to model the relationship between two variables: a dependent variable (the target - here, a house's sale price) and an independent variable (the predictor). The goal of simple linear regression is to find the best-fitting linear equation that describes the relationship between these variables.

1. **Purpose:** Simple linear regression is used to understand and quantify the relationship between two variables. It helps predict the value of the dependent variable based on the value of the independent variable.
2. **Assumptions:** The model assumes a linear relationship between the variables, constant variance of errors (homoscedasticity), and normally distributed residuals.

5. a- Identifying a highly correlated predictor

The target variable is price. Looking at correlation coefficients for all predictor variables to find the one with the highest correlation with price.

```
In [12]: kc_house_df.corr()["price"].sort_values(ascending=False)
```

```
Out[12]: price      1.000000
          sqft_living   0.608616
          sqft_above    0.538631
          bathrooms     0.480337
          sqft_patio    0.313789
          bedrooms      0.288954
          sqft_garage   0.263674
          sqft_basement 0.245005
          floors        0.180589
          yr_built      0.095796
          sqft_lot       0.086550
          Name: price, dtype: float64
```

Looking at correlation coefficients for all predictor variables to find the one with the highest correlation with price.

```
In [13]: most_correlated = "sqft_living"
```

5. b- Visualizing the predictor vs price

Since the price data points have a high concentration around the short range of square footage, a regression line is calculated to include into the scatter plot and make the graph more easy to read.

```
In [14]: # Calculating regression line
```

```
x = kc_house_df["sqft_living"].array
y = kc_house_df["price"].array

m, b = np.polyfit(x, y, 1)
```

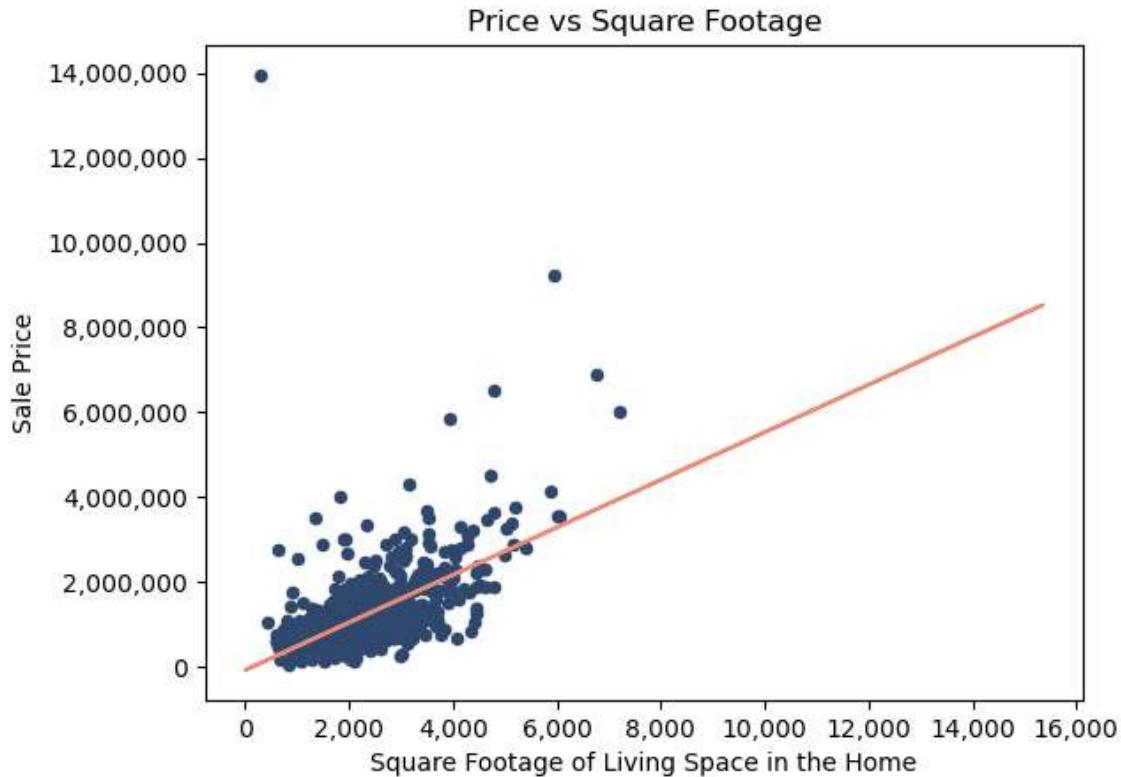
```
In [15]: ┌ # Plotting scatter plot
# Plotting a sample of 1000 data points, most_correlated vs. price
kc_house_df.sample(1000, random_state=1).plot.scatter(x=most_correlated, y=)

# Plotting regression line
ax = plt.plot(x, m*x+b, color='#ed8c72')

# Getting a reference to the current Axes object
ax = plt.gca()

# Formatting the x and y-axis tick labels to show actual numbers instead of scientific notation
ax.xaxis.set_major_formatter(ticker.FuncFormatter(lambda x, pos: '{:,.0f}'.format(x)))
ax.yaxis.set_major_formatter(ticker.FuncFormatter(lambda x, pos: '{:,.0f}'.format(x)))

# Defining titles and legend
plt.title('Price vs Square Footage')
plt.xlabel('Square Footage of Living Space in the Home')
plt.ylabel('Sale Price')
plt.savefig('images/sqft_living.png', dpi=300)
plt.show();
```



While outliers can be identified on the sale price, the variables of a house's price and square footage of living space seem to have a linear relationship with one another, confirming a simple linear regression model can be calculated.

5. c- Setting up variables for regression

Declaring y and X_baseline variables, where y is a Series containing price data and X_baseline is a DataFrame containing the column with the strongest correlation. As defined before, the latter is sqft_living.

```
In [16]: ┆ y = kc_house_df["price"]  
X_baseline = kc_house_df[[most_correlated]]
```

5. d- Creating and fitting simple linear regression

```
In [17]: ┆ baseline_model = sm.OLS(y, sm.add_constant(X_baseline))  
baseline_results = baseline_model.fit()
```

5. e- Evaluating and interpreting baseline regression results

In [18]: ┏ baseline_results.summary()

Out[18]: OLS Regression Results

Dep. Variable:	price	R-squared:	0.370			
Model:	OLS	Adj. R-squared:	0.370			
Method:	Least Squares	F-statistic:	1.771e+04			
Date:	Mon, 04 Sep 2023	Prob (F-statistic):	0.00			
Time:	09:05:41	Log-Likelihood:	-4.4847e+05			
No. Observations:	30111	AIC:	8.969e+05			
Df Residuals:	30109	BIC:	8.970e+05			
Df Model:	1					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	-7.559e+04	9798.815	-7.714	0.000	-9.48e+04	-5.64e+04
sqft_living	560.5134	4.211	133.096	0.000	552.259	568.768
Omnibus:	43384.874	Durbin-Watson:			1.862	
Prob(Omnibus):	0.000	Jarque-Bera (JB):			47187801.072	
Skew:	8.195	Prob(JB):			0.00	
Kurtosis:	196.242	Cond. No.			5.56e+03	

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 5.56e+03. This might indicate that there are strong multicollinearity or other numerical problems.

sqft_living was the attribute the most strongly correlated with price. Consequently, the model is describing the relationship between both variables.

1. **Adj. R-squared:** Overall, the model is not very significant and explains only 37% of the price's variance.
2. **Coef and p-values (P>|t|):** The coefficient for sqft_living is about \$560.5134. This means that for each additional square foot of living space in the home, the price would increase by a little over \$560. The p-value associated for the coefficient is equal to 0, indicating that the variable has a statistically significant impact on price.
3. **Constant coef:** The intercept is -\$75,590, meaning that a zero square foot living house would sell for -\$75,590 - which is not representative as a house would not sell for a negative value, and would not have 0 square foot.

In conclusion, this model suggests that the square footage of living space in a house is a statistically significant predictor of the price. For each additional square footage added to a house, the predicted price increases by around \$560.51. Nevertheless, the low adjusted r-squared value indicates that the model is not significant and would benefit from being improved.

In addition, other statistics available question the assumption of linear regression. These will be more detailed among the future models.

5. f- Calculating and interpreting the model's mean absolute error



```
In [19]: ┆ baseline_mae = mean_absolute_error(y, baseline_results.predict(sm.add_const  
baseline_mae
```

```
Out[19]: 396271.4457372302
```

The mean absolute error is a common metric used to evaluate the performance of a regression model. It quantifies the average absolute difference between the predicted values and the actual (true) values.

A lower mean absolute error indicates that the model's predictions are closer to the actual values. In the opposite, the MAE of around 396271.445 implies that, on average, the baseline model's predictions have an absolute difference of approximately \$396,271.445 from the actual target values. This confirms the poor performance of the initial baseline regression model's results.

6. Data Transformation

6. a- Visualizing the distribution of target variable price

```
In [20]: # Creating a figure with two subplots in one row and two columns
fig, axes = plt.subplots(1, 2, sharex=True, figsize=(26,13))

# Setting style
sns.set(style="white")

# First subplot: histplot and the mean price
sns.histplot(ax=axes[0], data=kc_house_df, x="price", bins=50, color="#2f497f")
mean_price = kc_house_df["price"].mean()
axes[0].axvline(x=mean_price, color="#ed8c72", linestyle='--', label='Mean')

# First subplot formatting
title_font = {'size': 24}
axes[0].set_title("Distribution of Price with Mean", fontdict=title_font)
axes[0].set_xlabel("Price", fontsize=20)
axes[0].set_ylabel("Frequency", fontsize=20)
axes[0].legend()

# Formatting the x-axis tick Labels to show actual numbers instead of scientific notation
axes[0].xaxis.set_major_formatter(ticker.FuncFormatter(lambda x, pos: '{:,}'))
axes[0].yaxis.set_major_formatter(ticker.FuncFormatter(lambda x, pos: '{:,}'))
axes[0].tick_params(labelsize=20)

# Second subplot: boxplot
sns.boxplot(ax=axes[1], data=kc_house_df, x="price", color="#2f496e")

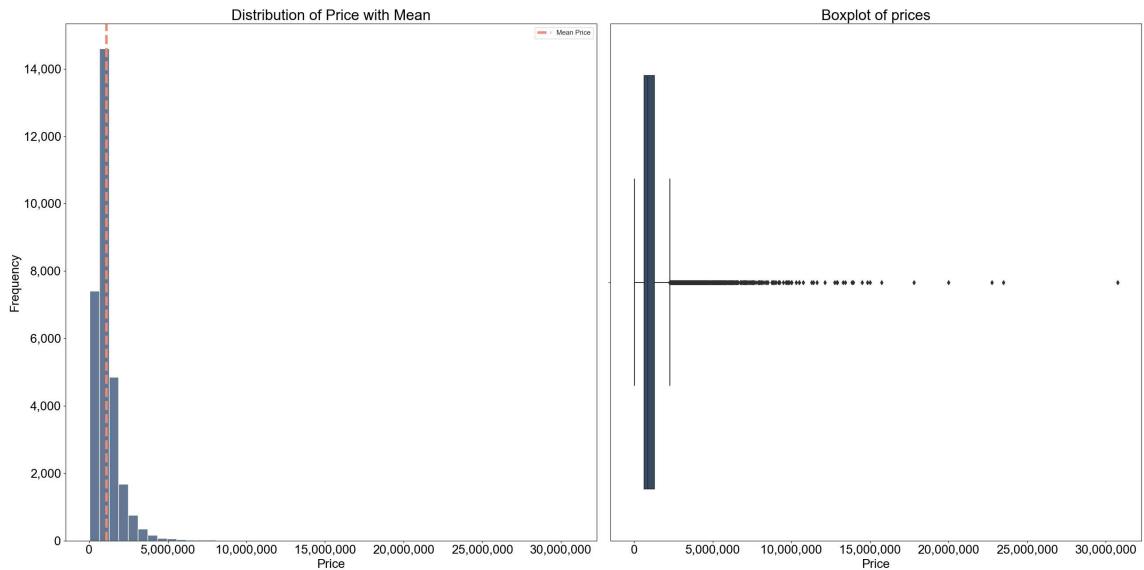
# Second subplot formatting
axes[1].set_title("Boxplot of prices", fontdict=title_font)
axes[1].set_xlabel("Price", fontsize=20)

# Formatting the x-axis tick Labels to show actual numbers instead of scientific notation
axes[1].xaxis.set_major_formatter(ticker.FuncFormatter(lambda x, pos: '{:,}'))
axes[1].yaxis.set_major_formatter(ticker.FuncFormatter(lambda x, pos: '{:,}'))
axes[1].tick_params(labelsize=20)

# Adjusting Layout
plt.tight_layout()

# Saving image
plt.savefig('images/price_distribution_before.png', dpi=300)

plt.show()
```



```
In [21]: median_price = kc_house_df["price"].median()
median_price = "${:,.0f}".format(median_price)
```

```
In [22]: mean_price = "${:,.0f}".format(mean_price)
```

```
In [23]: print(f"The highest number of houses are sold below the mean price of {mean
```

The highest number of houses are sold below the mean price of \$1,108,971. The mean is higher than the median (\$860,000) which indicates the distribution of the dataset is skewed to the right. There are less prices sold at higher prices, but the prices are so high compared to the most common sales prices that the mean overestimates the most common values.

The prices outliers will be identified and excluded to improve the model.

6. b- Identifying and excluding outliers

Sales prices were considered outliers when outside of the 5th percentile

The 5th percentile value was chosen as a standard percentile value to define outliers. The below cells display how they were calculated.

```
In [24]: min_price = kc_house_df["price"].min()
max_price = kc_house_df["price"].max()
```

```
In [25]: lower_limit = np.percentile(kc_house_df["price"], 0.5)
upper_limit = np.percentile(kc_house_df["price"], 99.5)
```

```
In [26]: ┏ ━ # Ensuring that the lower limit defined is not lower than the minimum price  
assert min_price < lower_limit
```

```
      # Ensuring that the upper limit defined is not higher than the maximum price  
assert max_price > upper_limit
```

```
In [27]: ┏ ━ # Now modifying the dataset excluding prices' outliers  
kc_house_df = kc_house_df[(kc_house_df["price"] > lower_limit) & (kc_house_
```

```
In [28]: ┏ ━ # Verifying the new minimum and maximum prices  
print(kc_house_df["price"].min())  
print(kc_house_df["price"].max())
```

```
131106.0  
5650000.0
```

6. c- Verifying the new distribution of target variable price

```
In [29]: # Creating a figure with two subplots in one row and two columns
fig, axes = plt.subplots(1, 2, sharex=True, figsize=(26,13))

# Setting style
sns.set(style="white")

# First subplot: histplot and the mean price
sns.histplot(ax=axes[0], data=kc_house_df, x="price", bins=50, color="#2f497f")
mean_price = kc_house_df["price"].mean()
axes[0].axvline(x=mean_price, color="#ed8c72", linestyle='--', label='Mean')

# First subplot formatting
title_font = {'size': 24}
axes[0].set_title("Distribution of Price with Mean", fontdict=title_font)
axes[0].set_xlabel("Price", fontsize=20)
axes[0].set_ylabel("Frequency", fontsize=20)
axes[0].legend()

# Formatting the x-axis tick Labels to show actual numbers instead of scientific notation
axes[0].xaxis.set_major_formatter(ticker.FuncFormatter(lambda x, pos: '{:,}'))
axes[0].yaxis.set_major_formatter(ticker.FuncFormatter(lambda x, pos: '{:,}'))
axes[0].tick_params(labelsize=20)

# Second subplot: boxplot
sns.boxplot(ax=axes[1], data=kc_house_df, x="price", color="#2f496e")

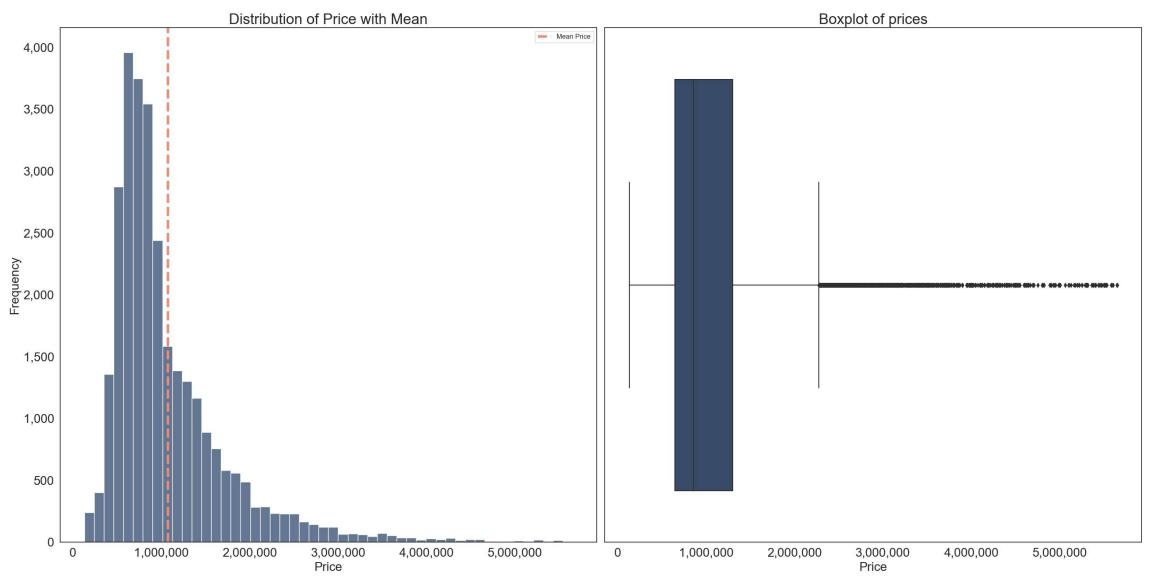
# Second subplot formatting
axes[1].set_title("Boxplot of prices", fontdict=title_font)
axes[1].set_xlabel("Price", fontsize=20)

# Formatting the x-axis tick Labels to show actual numbers instead of scientific notation
axes[1].xaxis.set_major_formatter(ticker.FuncFormatter(lambda x, pos: '{:,}'))
axes[1].tick_params(labelsize=20)

# Adjusting layout
plt.tight_layout()

# Saving image
plt.savefig('images/price_distribution_after.png', dpi=300)

plt.show()
```



6. d- Verifying the new correlation of square footage of living area and price

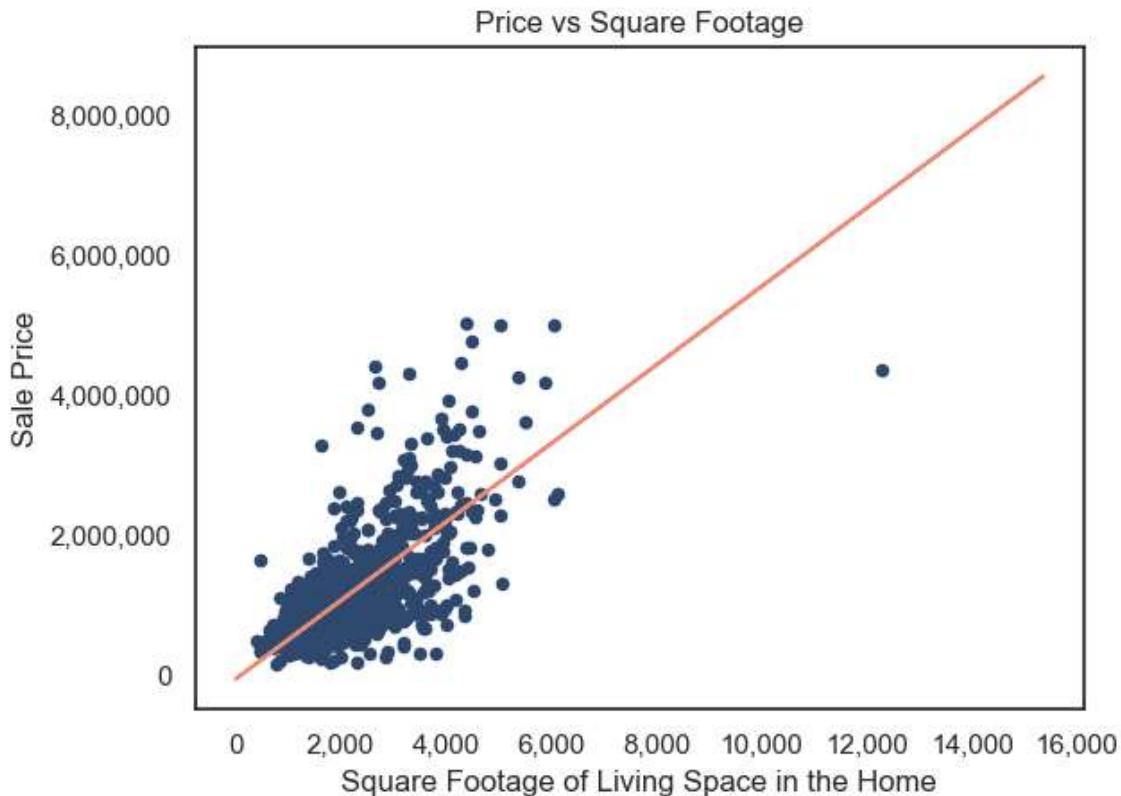
```
In [30]: # Plotting scatter plot
# Plotting a sample of 1000 data points, most_correlated vs. price
kc_house_df.sample(1000, random_state=1).plot.scatter(x=most_correlated, y=)

# Plotting regression Line
ax = plt.plot(x, m*x+b, color='#ed8c72')

# Getting a reference to the current Axes object
ax = plt.gca()

# Divide the tick labels by 1000 to display values in thousands
ax.xaxis.set_major_formatter(ticker.FuncFormatter(lambda x, pos: '{:,.0f}'.format(x/1000)))
ax.yaxis.set_major_formatter(ticker.FuncFormatter(lambda x, pos: '{:,.0f}'.format(x/1000)))

# Defining titles and Legend
plt.title('Price vs Square Footage')
plt.xlabel('Square Footage of Living Space in the Home')
plt.ylabel('Sale Price')
plt.savefig('images/sqft_living_post_outliers.png', dpi=300)
plt.show();
```



The distribution of the dataset seems now more equally distributed in the sense that outliers account for a lower portion of the data. Nevertheless, it still seems skewed and not to represent a normal distribution.

6. e- Inspecting the new statistics

In [31]: ► `# Inspecting statistics about the dataset
kc_house_df.describe()`

Out[31]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors
count	2.980700e+04	29807.000000	29807.000000	29807.000000	2.980700e+04	29807.000000
mean	1.076879e+06	3.412252	2.327893	2100.225316	1.645528e+04	1.543211
std	6.832233e+05	0.975721	0.870172	938.698138	5.872732e+04	0.567711
min	1.311060e+05	0.000000	0.000000	3.000000	4.020000e+02	1.000000
25%	6.500000e+05	3.000000	2.000000	1420.000000	4.827500e+03	1.000000
50%	8.600000e+05	3.000000	2.500000	1920.000000	7.448000e+03	1.500000
75%	1.300000e+06	4.000000	3.000000	2610.000000	1.050000e+04	2.000000
max	5.650000e+06	13.000000	10.500000	15360.000000	3.253932e+06	4.000000

7. Second Simple Linear Regression: Without Outliers

7. a- Setting up variables for regression

The dependent and independent variables are the same: y is the house's sale price and X_baseline, sqft_living.

In [32]: ► `# Declaring the variables once again for clarity
y = kc_house_df["price"]

new_X_baseline = kc_house_df[[most_correlated]]`

7. b- Creating and fitting simple linear regression

In [33]: ► `new_baseline_model = sm.OLS(y, sm.add_constant(new_X_baseline))
new_baseline_results = new_baseline_model.fit()`

7. c- Evaluating and interpreting baseline regression results

In [34]: ┏ new_baseline_results.summary()

Out[34]: OLS Regression Results

Dep. Variable:	price	R-squared:	0.428			
Model:	OLS	Adj. R-squared:	0.428			
Method:	Least Squares	F-statistic:	2.226e+04			
Date:	Mon, 04 Sep 2023	Prob (F-statistic):	0.00			
Time:	09:05:45	Log-Likelihood:	-4.3442e+05			
No. Observations:	29807	AIC:	8.689e+05			
Df Residuals:	29805	BIC:	8.689e+05			
Df Model:	1					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	7.733e+04	7337.847	10.539	0.000	6.29e+04	9.17e+04
sqft_living	475.9234	3.190	149.204	0.000	469.671	482.175
Omnibus:	12050.548	Durbin-Watson:		1.937		
Prob(Omnibus):	0.000	Jarque-Bera (JB):		98988.100		
Skew:	1.734	Prob(JB):		0.00		
Kurtosis:	11.227	Cond. No.		5.64e+03		

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.64e+03. This might indicate that there are strong multicollinearity or other numerical problems.

1. **Adjusted R-squared:** Removing the outliers highly improved the model which now explains 42.8% of the price's variance that can be explained by square foot living.
2. **Coef and p-values (P<|t|):** Because outliers were removed, the coefficient decreased to 475.9234, indicating that the selling price would increase by a lower amount than previously recorded: \$475.92. The p-value of 0 confirms the square footage of living space in the house has a significant impact on the price.
3. **Constant coef:** The intercept value is now 77330, indicating the estimated price when the square footage of living space in the house is at 0 is \$77,330 - which is much more realistic than the previously recorded negative number.
4. **F-statistic and Prob (F-statistic):** A high F-statistic combined with a low p-value suggests that the model's overall fit is statistically significant. p-values were at 0 on both models, but F-statistic increased compared to the previous simple linear regression: 22,260 vs 17,710.

This confirms the model improved from the initial one.

7. d- Calculating and interpreting model's mean absolute error

In [35]: ► new_baseline_mae = mean_absolute_error(y, new_baseline_results.predict(sm.a
baseline_mae, new_baseline_mae

```
# new_baseline_mae - baseline_mae  
baseline_mae, new_baseline_mae
```

Out[35]: (396271.4457372302, 355912.7383624806)

The new MAE is 40,358 less than the previous one, indicating the model's predictions are closer to the actual values. The new baseline model's predictions have an absolute difference of approximately \$355,912 from the actual target values.

The range of sales prices is extremely large: it goes from \$131,106 to \$5,650,000. This would make this MAE acceptable in a real-world context. Nevertheless, this model takes only one variable into account. In addition scaling the data could help improve prediction accuracy.

8. Multiple Linear Regression - Categorical Features

8. a- Inspecting categorical features

```
In [36]: # Inspecting each categorical feature to understand their attributes
categoricals = kc_house_df.select_dtypes(include=[object])

for col in categoricals:
    print(kc_house_df[col].value_counts(), "\n")
```

NO 24646
YES 5161
Name: nuisance, dtype: int64

Average 18301
Good 7988
Very Good 3237
Fair 224
Poor 57
Name: condition, dtype: int64

7 Average 11615
8 Good 9352
9 Better 3791
6 Low Average 2807
10 Very Good 1344
5 Fair 381
11 Excellent 362
12 Luxury 92
4 Low 45
13 Mansion 9
3 Poor 8
2 Substandard 1
Name: grade, dtype: int64

Gas 20382
Electricity 6389
Oil 2865
Gas/Solar 90
Electricity/Solar 58
Other 19
Oil/Solar 4
Name: heat_source, dtype: int64

PUBLIC 25501
PRIVATE 4299
PRIVATE RESTRICTED 4
PUBLIC RESTRICTED 3
Name: sewer_system, dtype: int64

There are 5 columns with categorical data:

- nuisance
- condition
- grade
- heat_source
- sewer_system

Nuisance's response is binary: yes or no. Condition and grade provide a rating system from poor to luxury. Rating includes some specific to the house type, i.e. mansion.

Heat source and sewer system's responses refer to options specific to each category, i.e. gas, electricity as heat sources in the house.

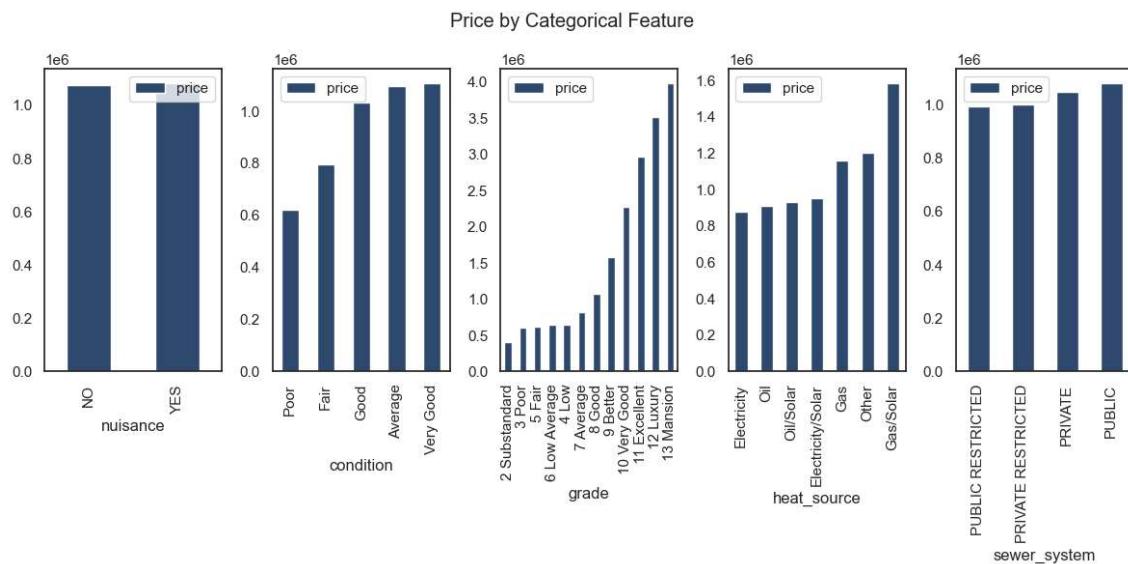
8. b- Identifying a promising predictor

```
In [37]: ┌─ categorical_features = categoricals.columns
```

```
In [38]: ┌─ categorical_features
```

```
Out[38]: Index(['nuisance', 'condition', 'grade', 'heat_source', 'sewer_system'],  
              dtype='object')
```

```
In [39]: ┌─ fig, axes = plt.subplots(ncols=len(categorical_features), figsize=(12, 6))  
          for index, feature in enumerate(categoricals):  
              kc_house_df.groupby(feature).mean().sort_values("price").plot.bar(y="pr  
          plt.suptitle("Price by Categorical Feature")  
          plt.tight_layout()  
          plt.show()
```



Most categorical features seem to correlate with price except for nuisance.

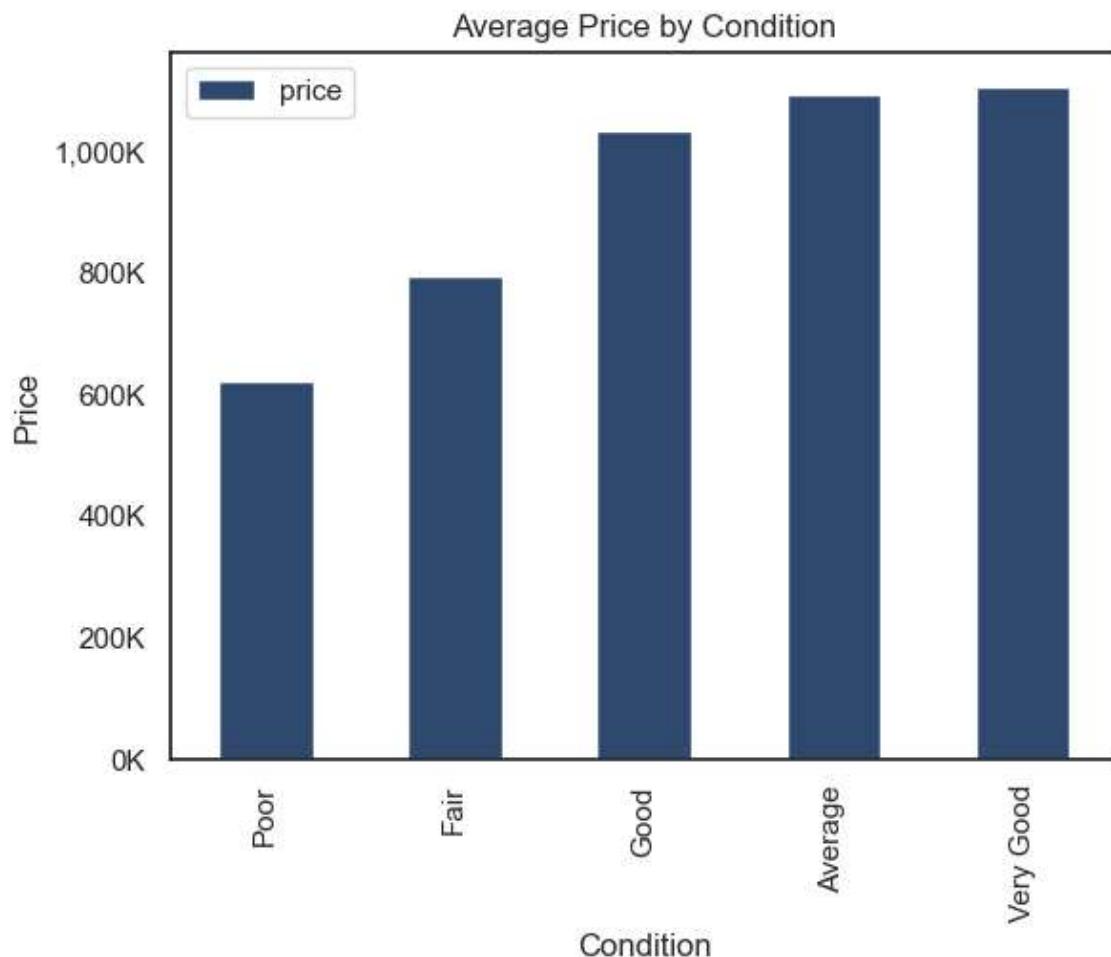
We will start with the overall representation: condition, as even though grade gives an idea, reaching a mansion by renovation is less likely.

```
In [40]: cat_feature = "condition"
```

```
In [41]: ax = kc_house_df.groupby(cat_feature)[ "price" ].mean().reset_index().sort_values(
```

Customize the y-axis tick Labels to show values in thousands
ax.yaxis.set_major_formatter(ticker.FuncFormatter(lambda y, pos: f'{y/1000}':

```
plt.title("Average Price by Condition")  
plt.xlabel('Condition')  
plt.ylabel('Price')  
plt.savefig('images/price_by_condition.png', dpi=500)  
plt.show()
```



8. c- Setting up variables for regression

Creating a variable `X_iterated` : a DataFrame containing the column with the strongest correlation and the selected categorical feature.

```
In [42]: X_iterated = kc_house_df[[most_correlated, cat_feature]]  
X_iterated
```

Out[42]:

	sqft_living	condition
0	1180	Good
1	2770	Average
2	2880	Average
3	2160	Average
4	1120	Average
...
30150	1910	Good
30151	2020	Average
30152	1620	Average
30153	2570	Average
30154	1200	Average

29807 rows × 2 columns

8. d- Preprocessing categorical variable with one-hot encoding

```
In [43]: X_iterated = pd.get_dummies(X_iterated, columns=[cat_feature])  
X_iterated
```

Out[43]:

	sqft_living	condition_Average	condition_Fair	condition_Good	condition_Poor	condit
0	1180	0	0	1	0	
1	2770	1	0	0	0	
2	2880	1	0	0	0	
3	2160	1	0	0	0	
4	1120	1	0	0	0	
...
30150	1910	0	0	1	0	
30151	2020	1	0	0	0	
30152	1620	1	0	0	0	
30153	2570	1	0	0	0	
30154	1200	1	0	0	0	

29807 rows × 6 columns

```
In [44]: ┏ ━ # Dropping "condition_Poor" as a baseline, because is the lowest ranking gi
X_iterated.drop("condition_Poor", axis=1, inplace=True)
X_iterated
```

Out[44]:

	sqft_living	condition_Average	condition_Fair	condition_Good	condition_Very Good
0	1180	0	0	1	0
1	2770	1	0	0	0
2	2880	1	0	0	0
3	2160	1	0	0	0
4	1120	1	0	0	0
...
30150	1910	0	0	1	0
30151	2020	1	0	0	0
30152	1620	1	0	0	0
30153	2570	1	0	0	0
30154	1200	1	0	0	0

29807 rows × 5 columns

8. e- Building a multiple linear regression model

Using the y variable from our previous model and X_iterated , building a model called iterated_model and a regression results object called iterated_results .

```
In [45]: ┏ ━ iterated_model = sm.OLS(y, sm.add_constant(X_iterated))
iterated_results = iterated_model.fit()
```

```
In [46]: ┏ ━ coefficients = iterated_results.params[1:]
intercept = iterated_results.params[0]
print(coefficients)
print(intercept)
```

```
sqft_living          477.250792
condition_Average   51872.299370
condition_Fair       62919.975143
condition_Good      69085.793874
condition_Very Good 135739.940053
dtype: float64
8967.235833342129
```

8. f- Evaluating and interpreting regression results

In [47]: `iterated_results.summary()`

Out[47]: OLS Regression Results

Dep. Variable:	price	R-squared:	0.429			
Model:	OLS	Adj. R-squared:	0.429			
Method:	Least Squares	F-statistic:	4478.			
Date:	Mon, 04 Sep 2023	Prob (F-statistic):	0.00			
Time:	09:05:46	Log-Likelihood:	-4.3439e+05			
No. Observations:	29807	AIC:	8.688e+05			
Df Residuals:	29801	BIC:	8.688e+05			
Df Model:	5					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	8967.2358	6.85e+04	0.131	0.896	-1.25e+05	1.43e+05
sqft_living	477.2508	3.204	148.959	0.000	470.971	483.531
condition_Average	5.187e+04	6.86e+04	0.757	0.449	-8.25e+04	1.86e+05
condition_Fair	6.292e+04	7.66e+04	0.821	0.411	-8.72e+04	2.13e+05
condition_Good	6.909e+04	6.87e+04	1.006	0.314	-6.55e+04	2.04e+05
condition_Very Good	1.357e+05	6.9e+04	1.966	0.049	440.895	2.71e+05
Omnibus:	12119.347	Durbin-Watson:	1.935			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	100578.330			
Skew:	1.743	Prob(JB):	0.00			
Kurtosis:	11.297	Cond. No.	1.18e+05			

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.18e+05. This might indicate that there are strong multicollinearity or other numerical problems.

This multiple linear regression analysis is aimed at predicting the "price" variable based on a set of independent variables: square footage of living are in the house, as well as the various rankings.

1. **R-squared and Adjusted R-squared:** The adjusted R-squared is 0.429. This is the most significant model thus far. It explains 42.9% of the price's variance.

2. **F-statistic and Prob (F-statistic):** The F-statistic of 4478 with a very low probability (Prob (F-statistic) = 0) indicates that at least one independent variable in the model has a significant relationship with price.
3. **Coefficients (coef) and P-values (P>|t|):** Each coefficient represents the change in the dependent variable for a one-unit change in the corresponding independent variable, while holding other variables constant. For example:

- "sqft_living" coefficient: An increase of 1 square foot in living area is associated with an increase of \$477.25 in price - which is consistent with the previous simple linear regression (475.92).
- "condition_Average" coefficient: Properties with an "Average" condition are associated with an increase of \$51,870 in price, but the p-value is not significant ($p > 0.05$).
- "condition_Fair" coefficient: Properties with a "Fair" condition are associated with an increase of \$62,920 in price, but the p-value is not significant.
- "condition_Good" coefficient: Properties with a "Good" condition are associated with an increase of \$69,090 in price, but the p-value is not significant.
- "condition_Very Good" coefficient: Properties with a "Very Good" condition are associated with an increase of \$135,700 in price, and the p-value (0.049) indicates that this increase is statistically significant at a 5% significance level.

Limits

1. **Jarque-Bera, Skew, Kurtosis:** This high Jarque-Bera value combined with a low p-value indicate that the residuals are not normally distributed and there might be issues with the model's assumptions. The skewness (1.73) suggests that the residuals are positively skewed, and the kurtosis value (11.297) indicates heavy tails in the distribution.
2. **Durbin-Watson:** The Durbin-Watson statistic of 1.935 suggests that there may be some positive autocorrelation among the residuals.
3. **Cond. No.:** The condition number provides information about multicollinearity. A high value suggests that there might be multicollinearity among the independent variables.

In conclusion, this regression analysis on categorical features indicates that the model has some

8. g- Calculating and interpreting model's mean absolute error

```
In [48]: ┏ iterated_mae = mean_absolute_error(y, iterated_results.predict(sm.add_const
baseline_mae, new_baseline_mae, iterated_mae)
```

```
Out[48]: (396271.4457372302, 355912.7383624806, 355653.4317614562)
```

```
In [49]: ┏ baseline_results.rsquared_adj, new_baseline_results.rsquared_adj, iterated_
```

```
Out[49]: (0.37039218894923187, 0.4275443480727944, 0.4288815982113814)
```

This model is predicting price based on square footage of living area and the house's condition returned the highest adjusted r-squared. The Mean Absolute Error also indicates the absolute differences between predicted and actual values are the smallest of all models.

9. Multiple Linear Regression - Numeric Features

9. a- Setting up variables for regression

```
In [50]: # Setting categorical column names as the variable categorical_features
categorical_features = categoricals.columns
```



```
In [51]: # Dropping the dependent variable - the column price, from the dataset
house_preds = kc_house_df.drop("price", axis=1)

# Defining the independent variables: removing categorical columns from house_preds
house_preds = house_preds.drop(categorical_features, axis=1)

# Defining the target variable: price
house_target = kc_house_df["price"]

# Inspecting the new data
house_preds.head()
```

Out[51]:

	bedrooms	bathrooms	sqft_living	sqft_lot	floors	sqft_above	sqft_basement	sqft_garag
0	4	1.0	1180	7140	1.0	1180	0	1
1	5	2.5	2770	6703	1.0	1570	1570	1
2	6	2.0	2880	6156	1.0	1580	1580	1
3	3	3.0	2160	1400	2.0	1090	1070	20
4	2	2.0	1120	758	2.0	1120	550	55



We are adding another predictor call const. The value of const is always 1. It is the intercept term for our regression equation. By multiplying beta to 1, the constant, we will get the raw beta.

Now building our model.

9. b- Building a multiple linear regression model

```
In [52]: # Adding the constant predictor
predictors = sm.add_constant(house_preds)
```



```
In [53]: numeric_results = sm.OLS(house_target, predictors).fit()
```

9. c- Evaluating and interpreting regression results

In [54]: ┏ numeric_results.summary()

Out[54]: OLS Regression Results

Dep. Variable:	price	R-squared:	0.464			
Model:	OLS	Adj. R-squared:	0.464			
Method:	Least Squares	F-statistic:	2579.			
Date:	Mon, 04 Sep 2023	Prob (F-statistic):	0.00			
Time:	09:05:46	Log-Likelihood:	-4.3344e+05			
No. Observations:	29807	AIC:	8.669e+05			
Df Residuals:	29796	BIC:	8.670e+05			
Df Model:	10					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	6.655e+06	2.51e+05	26.515	0.000	6.16e+06	7.15e+06
bedrooms	-1.091e+05	4041.609	-26.986	0.000	-1.17e+05	-1.01e+05
bathrooms	1.162e+05	5885.899	19.734	0.000	1.05e+05	1.28e+05
sqft_living	303.2609	13.609	22.283	0.000	276.586	329.936
sqft_lot	0.0019	0.051	0.037	0.970	-0.098	0.101
floors	3.9e+04	7124.309	5.474	0.000	2.5e+04	5.3e+04
sqft_above	202.6998	13.896	14.587	0.000	175.463	229.936
sqft_basement	101.1899	10.177	9.943	0.000	81.243	121.137
sqft_garage	-51.8612	14.017	-3.700	0.000	-79.335	-24.387
sqft_patio	159.9047	13.283	12.039	0.000	133.870	185.939
yr_built	-3343.4651	130.121	-25.695	0.000	-3598.507	-3088.423
Omnibus:	11889.713	Durbin-Watson:	1.931			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	105797.339			
Skew:	1.679	Prob(JB):	0.00			
Kurtosis:	11.597	Cond. No.	5.28e+06			

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 5.28e+06. This might indicate that there are strong multicollinearity or other numerical problems.

This mutiple linear regression analysis is aimed at predicting the "price" variable based on the numeric features: the independent variables.

1. **Adjusted R-squared:** The R-squared and adjusted R-squared values of 0.464 indicate that about 46.4% of the variance in ("price") is explained by the independent variables in the model, and that the model's significance is the highest.
2. **F-statistic and Prob (F-statistic):** The F-statistic of 2578 with a very low probability (Prob (F-statistic) = 0) indicates that at least one of the numeric features has a significant relationship with the sales price.
3. **Coefficients (coef) and P-values (P>|t|):** Each coefficient represents the change in the dependent variable for a one-unit change in the corresponding independent variable, while holding other variables constant. And the p-values associated with them test the null hypothesis that the varibale has a statistically significant impact on the dependent variable. Unless specified, all p-values confirmed the statistically significant impact of each feature on price.
 - "sqft_lot":an increase of one square foot in lot area is associated with an increase of \$0.0019 on the sale price. Not only this value is very low, but the p-value is greater than 0.05 which indicates this feature does not have a significant impact on the price. For better results, this feature should be dropped to improve the model's results.
 - "bedrooms", sqft_garage and "yr_built" return negative coefficients which would indicate that increase the number of bedrooms, square feet for a garage and the higher the year built would negatively impact the sales price.

Limits

1. **Jarque-Bera, Skew, Kurtosis:** Again, the high Jarque-Bera value (105797.34) associated with a low p-value (0) indicates that the residuals are not normally distributed. The skewness value (1.68) confirms the distribution is highly skewed and a high Kurtosis (11.58) validates there are heavy tails. This questions the model assumptions that the distribution is not normally distributed.
2. **Durbin-Watson:** The Durbin-Watson statistic of 1.93, being close to 2, indicates a low risk of positive autocorrelation among the residuals, but suggests there might be some.
3. **Cond. No.:** The high condition number (5,280,000) suggests that there may be multicollinearity among the independent variables.

9. d- Calculating and interpreting model's mean absolute error

In [55]: ┏ numeric_mae = mean_absolute_error(house_target, numeric_results.predict(prep_baseline_mae, iterated_mae, numeric_mae))

Out[55]: (396271.4457372302, 355653.4317614562, 340763.44152917946)

In [56]: ┏ baseline_results.rsquared_adj, iterated_results.rsquared_adj, numeric_resul

Out[56]: (0.37039218894923187, 0.4288815982113814, 0.46381145072171526)

This model is predicting price based on the numeric features of the dataset. It returned the highest adjusted r-squared. The Mean Absolute Error also indicates the absolute differences between predicted and actual values are the smallest of all models.

10. Multiple Linear Regression - New Model - Numeric Features

10. a- Setting up variables for regression

"sqft_lot" returned a low coefficient with a high p-value, suggesting this feature does not have a significant impact on the house's sales price. It is being dropped to improve the model.

The target remains price, so does not need to be defined again.

```
In [57]: # Storing feature to drop in the variable "to_drop"
to_drop = ["sqft_lot"]

# Dropping them from the predictors
new_house_preds = house_preds.drop(to_drop, axis=1)
```

```
In [58]: # Inspecting the new data
new_house_preds.head()
```

Out[58]:

	bedrooms	bathrooms	sqft_living	floors	sqft_above	sqft_basement	sqft_garage	sqft_p:
0	4	1.0	1180	1.0	1180	0	0	
1	5	2.5	2770	1.0	1570	1570	0	
2	6	2.0	2880	1.0	1580	1580	0	
3	3	3.0	2160	2.0	1090	1070	200	
4	2	2.0	1120	2.0	1120	550	550	

Now building our model.

10. b- Building the new multiple linear regression model

```
In [59]: # Adding the constant predictor
new_predictors = sm.add_constant(new_house_preds)
```

```
In [60]: new_numeric_results = sm.OLS(house_target, new_predictors).fit()
```

10. c- Evaluating and interpreting multiple linear regression model results

In [61]: ► new_numeric_results.summary()

Out[61]: OLS Regression Results

Dep. Variable:	price	R-squared:	0.464			
Model:	OLS	Adj. R-squared:	0.464			
Method:	Least Squares	F-statistic:	2866.			
Date:	Mon, 04 Sep 2023	Prob (F-statistic):	0.00			
Time:	09:05:46	Log-Likelihood:	-4.3344e+05			
No. Observations:	29807	AIC:	8.669e+05			
Df Residuals:	29797	BIC:	8.670e+05			
Df Model:	9					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	6.655e+06	2.51e+05	26.518	0.000	6.16e+06	7.15e+06
bedrooms	-1.091e+05	4029.747	-27.069	0.000	-1.17e+05	-1.01e+05
bathrooms	1.161e+05	5880.475	19.750	0.000	1.05e+05	1.28e+05
sqft_living	303.2714	13.606	22.289	0.000	276.602	329.940
floors	3.898e+04	7101.652	5.489	0.000	2.51e+04	5.29e+04
sqft_above	202.7210	13.884	14.601	0.000	175.508	229.934
sqft_basement	101.1858	10.176	9.944	0.000	81.240	121.131
sqft_garage	-51.8570	14.016	-3.700	0.000	-79.330	-24.384
sqft_patio	159.9648	13.184	12.134	0.000	134.124	185.805
yr_built	-3343.4899	130.117	-25.696	0.000	-3598.524	-3088.456
Omnibus:	11890.126	Durbin-Watson:	1.931			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	105806.429			
Skew:	1.679	Prob(JB):	0.00			
Kurtosis:	11.598	Cond. No.	3.13e+05			

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 3.13e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [62]: # Reading all adjusted r-squared
```

```
baseline_results.rsquared_adj, new_baseline_results.rsquared_adj, iterated_
```

```
Out[62]: (0.37039218894923187,  
          0.4275443480727944,  
          0.4288815982113814,  
          0.46381145072171526,  
          0.46382942067496613)
```

The model's results are very similar to the initial multilinear regression model. A few statistics show improvement of the model:

- Adjusted R-squared:** the new adjusted r-squared shows a small improvement with a value of 0.46382942067496613 as opposed to 0.46381145072171526. 46.4% of the model's explanatory power is a little stronger than the previous one.
 - F-statistic and Prob (F-statistic):** the F-statistic value increased to 2866 from 2579 with a p-value remaining at 0. This indicates that at least one independent variable in the model has an even stronger impact on the price than the previous model.
 - Coefficients and P-values (P>|t|):** each coefficient represents the change in price for a one-unit change in the corresponding numeric, independent variable, while holding other variables constant. The p-values associated with them test the null hypothesis that the variable has a statistically significant impact on the dependent variable. All p-values were at 0, confirming the statistically significant impact of each feature on price.
 - "bathrooms": adding one bathroom to the house is associated with an increase of \$116,200 on the sale price
 - "sqft_living": an increase of one square foot in living area is associated with an increase of \$303.26 on the sale price
 - "floors": adding one floor to a house is associated with an increase of \$33,900 on the sale price
 - "sqft_above": an increase of one square foot in the above area is associated with an increase of \$202.70 on the sale price
 - "sqft_basement": an increase of one square foot in the basement area is associated with an increase of \$101.19 on the sale price
 - "sqft_patio": an increase of one square foot in the patio is associated with an increase of \$159.91 on the sale price
 - Condition no.:** although still high (313,000), this is a considerable improvement compared to the 5,280,000 recorded on the previous model. The risk of multicollinearity among the independent variables is now lower, improving the model's validity.

Limits

- 1. Jarque-Bera, and Kurtosis:** both statistics slightly increased on this new model, suggesting that removing "sqft_lot" worsened the distribution of the dataset, which was already not normal. The slightly higher Kurtosis (11.598 vs 11.597) indicates slightly heavier tails.

In summary, this final regression analysis indicates that the model has the most explanatory power among all models. The features having the most impact on sales price are bathrooms, sqft_living, floors, sqft_above, sqft_garage and sqft_patio. On the other hand, an increase in bedrooms sqft_garage and yr_built seem to lower the sale price. This may suggest that in order to increase the square footage of living area, the garage could be reduced or a bedroom could be turned into a larger living room to sell a house at a higher price.

A deeper analysis by house type should be conducted to confirm this hypothesis.

Nevertheless, there may be issues with the main assumptions of a multiple linear regression

10. d- Calculating and interpreting model's mean absolute error

```
In [63]: ┏ new_numeric_results.rsquared_adj - numeric_results.rsquared_adj
```

```
Out[63]: 1.7969953250873694e-05
```

```
In [64]: ┏ new_numeric_mae = mean_absolute_error(house_target, new_numeric_results.pre_baseline_mae, new_baseline_mae, iterated_mae, numeric_mae, new_numeric_mae)
```

```
Out[64]: (396271.4457372302,  
 355912.7383624806,  
 355653.4317614562,  
 340763.44152917946,  
 340762.8513420028)
```

```
In [65]: ┏ # Storing adjusted r-squared values from each model into variables  
baseline_rsq_adj = baseline_results.rsquared_adj  
new_baseline_rsq_adj = new_baseline_results.rsquared_adj  
iterated_rsq_adj = iterated_results.rsquared_adj  
numeric_rsq_adj = numeric_results.rsquared_adj  
new_numeric_rsq_adj = new_numeric_results.rsquared_adj
```

The Mean Absolute Error is now the lowest reached, indicating that the predicted values are the closest they have been from the actual values.

11. Summary: Visualizing Regression Models Results

```
In [66]: ┏ # transforming list to array to multiply adjusted r-squared values by 100.  
for_bar_data = [baseline_rsq_adj, new_baseline_rsq_adj, iterated_rsq_adj, n  
bar_data = np.array(for_bar_data)  
bar_data = bar_data * 100  
print(bar_data)
```

```
[37.03921889 42.75443481 42.88815982 46.38114507 46.38294207]
```

In [67]: ►

```
line_data = [baseline_mae, new_baseline_mae, iterated_mae, numeric_mae, new_x_values = ['Model 1', 'Model 2', 'Model 3', 'Model 4', 'Model 5']
```

```
In [68]: ┌ # Create a figure and primary axis
sns.set(style="white")

fig, ax1 = plt.subplots(figsize=(10, 6))

# Bar plot on the primary axis
sns.barplot(x=x_values, y=bar_data, ax=ax1, color="#2f496e")
# Set y-axis label for the bar plot
ax1.set_ylabel('Adjusted R-squared', color="#2f496e")
# Set y-axis limits for the primary axis
ax1.set_ylim(30, max(bar_data) + 2)

# Add data labels for each bar
for i, value in enumerate(bar_data):
    ax1.text(i, value + 0.5, f'{value:.4f}', ha='center', va='bottom', color='black')

ax2 = ax1.twinx()

# Line plot on the secondary axis
sns.lineplot(x=x_values, y=line_data, ax=ax2, color="#ed8c72")
ax2.set_ylabel('Mean Absolute Error', color="#ed8c72")

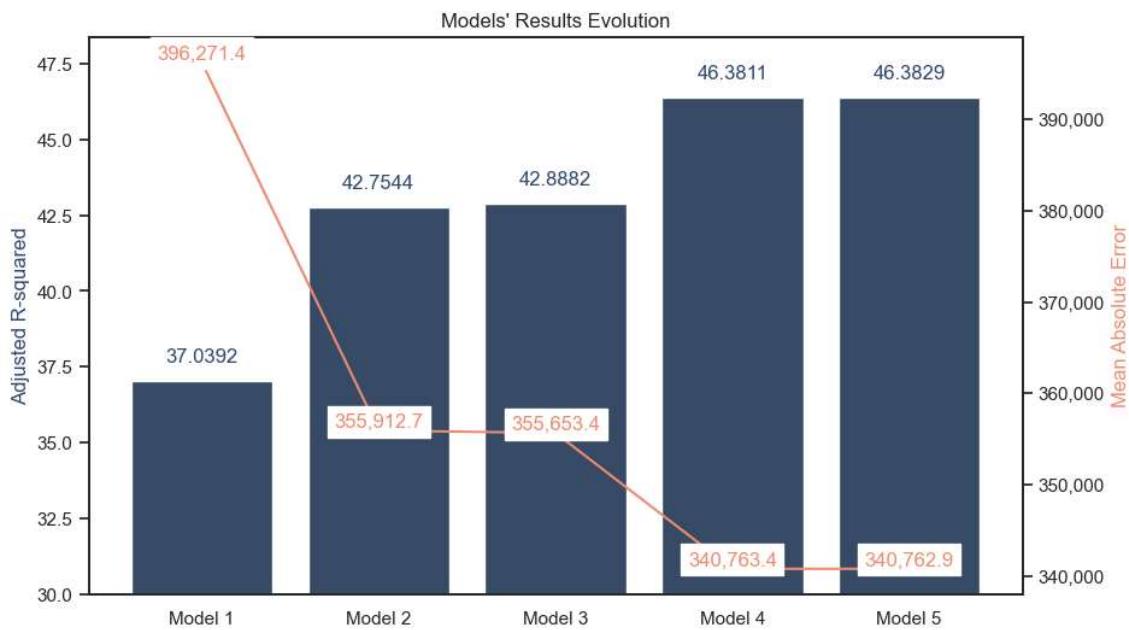
# Format y-axis of the secondary axis with thousand separators
ax2.yaxis.set_major_formatter(ticker.FuncFormatter(lambda x, pos: '{:, .0f}'))

# Add data labels for each bar
for i, value in enumerate(line_data):
    ax2.text(i, value + 0.1, f'{value:.1f}', ha='center', va='bottom', color='black')

plt.xlabel('Models')
plt.title("Models' Results Evolution" )

# Save the plot in images folder
plt.savefig('images/models_results.png', dpi=300)

plt.show()
```



12. Recommendations

The last model returned the best results to draw the below conclusions. As a real-estate, to improve your house's sale price:

1. Aim for a very good condition

Properties with a "Very Good" condition are associated with an increase of \$135,700 in price

2. Focus on the living area

Focus on increasing the living area, by maybe reducing the space of your garage, basement, or patio.

An increase of one square foot in living area is associated with an increase of \$303.26 on the sale price. If you increase your living room by 100 square feet, this represents a predicted increase of \$30,260.

3. Add one bathroom

If costs and water inlet permit it.

Adding one bathroom to the house is associated with an increase of \$116,200 on the sale price

4. Add one floor

If your house has high ceilings. Consider a mezzanine if height is limited.

Adding one floor to a house is associated with an increase of \$33,900 on the sale price

13. Limits and Next Steps

Limits

The last multiple linear regression results confirmed the findings of each model:

1. The dataset does not have a normal distribution
2. There is a rather risk of multicollinearity among the independent variables

Both findings present a risk on the model's assumptions and may question its validity.

Future Steps

The dataset could be transformed into a normal distribution in order to ensure a better fit. In addition, scaling the data could improve predictions, and reducing the Mean Absolute Error.

To address the risk of multicollinearity, some variables could be removed or grouped together.

**** Appendix ****

Second simple linear regression after removing outliers using Interquartile Range method

The Interquartile Range method to exclude outliers is used for normal distributions

In order to reach a similar Adjusted R-squared than by removing the 5th percentile, values outside of 8 times the IQR needed to be excluded as opposed to the usual values outside 1.5 times the IQR. This is because the dataset is not normally distributed. As a consequence, this method was not kept.

A. a- Identifying and excluding outliers

In [69]: ► *# Interquartile range method: 1. Calculating the Interquartile range*
q1 = kc_house_df["price"].quantile(0.25)
q3 = kc_house_df["price"].quantile(0.75)
IQR = q3 - q1

In [70]: ► *# 2. Define outlier thresholds. Typically, values outside of 1.5 times the*
lower_bound = q1 - 8 * IQR
upper_bound = q3 + 8 * IQR

In [71]: ► *# 3. Creating a mask to identify rows with prices that fall outside the def*
outliers_mask = (kc_house_df["price"] < lower_bound) | (kc_house_df["price"]

In [72]: ► *# 4. Extracting outliers*
not_outliers = kc_house_df[~outliers_mask]

A. b- Setting Up Variables for Regression (IQR method)

```
In [73]: ┏━ y = not_outliers["price"]
          X_baseline_IQR = not_outliers[["sqft_living"]]
```

A. c- Creating and Fitting Simple Linear Regression (IQR method)

```
In [74]: ┏━ baseline_model_IQR = sm.OLS(y, sm.add_constant(X_baseline_IQR))
          baseline_results_IQR = baseline_model_IQR.fit()
```

A. d- Baseline Model Results (IQR method)

```
In [75]: ┏━ baseline_results_IQR.summary()
```

Out[75]: OLS Regression Results

Dep. Variable:	price	R-squared:	0.428			
Model:	OLS	Adj. R-squared:	0.428			
Method:	Least Squares	F-statistic:	2.226e+04			
Date:	Mon, 04 Sep 2023	Prob (F-statistic):	0.00			
Time:	09:05:47	Log-Likelihood:	-4.3442e+05			
No. Observations:	29807	AIC:	8.689e+05			
Df Residuals:	29805	BIC:	8.689e+05			
Df Model:	1					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	7.733e+04	7337.847	10.539	0.000	6.29e+04	9.17e+04
sqft_living	475.9234	3.190	149.204	0.000	469.671	482.175
Omnibus:	12050.548	Durbin-Watson:	1.937			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	98988.100			
Skew:	1.734	Prob(JB):	0.00			
Kurtosis:	11.227	Cond. No.	5.64e+03			

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 5.64e+03. This might indicate that there are strong multicollinearity or other numerical problems.

